

```
In [1]: import pandas as pd
import numpy as np
import scipy.stats
```

C:\Users\hp\anaconda3\lib\site-packages\pandas\core\computation\expression
s.py:20: UserWarning: Pandas requires version '2.7.3' or newer of 'numexp
r' (version '2.7.1' currently installed).
from pandas.core.computation.check import NUMEXPR_INSTALLED

Test of Independence

$H_0: B_1 = 0$ (no relationship between x and y)

$H_a: B_1 \neq 0$ (there's a relationship)

- 1) Get F_0 from ANOVA table
 - 2) Get F_c from F table $F_c = F(\alpha, n - 2)$ where α is the level of significance
 - 3) Compare F_0 and F_c , if $F_0 > F_c$ we reject the null hypothesis which means that there is a relationship between the independent variable and the response
-

Interval Estimation

$$\hat{B}_0 \sim N(B_0, \sigma^2 \left(\frac{1}{n} + \frac{\bar{x}}{SXX} \right))$$

When σ^2 is known, $B_0: \hat{B}_0 \pm Z_{\frac{\alpha}{2}} \sqrt{\sigma^2 \left(\frac{1}{n} + \frac{\bar{x}}{SXX} \right)}$

When it's not, $B_0: \hat{B}_0 \pm t_{\frac{\alpha}{2}, n-2} \sqrt{MSE \left(\frac{1}{n} + \frac{\bar{x}}{SXX} \right)}$

$$\hat{B}_1 \sim N(B_1, \frac{\sigma^2}{SXX})$$

When σ^2 is known, $B_1: \hat{B}_1 \pm Z_{\frac{\alpha}{2}} \sqrt{\frac{\sigma^2}{SXX}}$

When it's not, $B_1: \hat{B}_1 \pm t_{\frac{\alpha}{2}, n-2} \sqrt{\frac{MSE}{SXX}}$

```
In [2]: df = pd.read_csv('income.data.csv')
```

```
In [3]: df.drop(columns='Unnamed: 0', inplace=True)
```

In [4]: `df.head()`

Out[4]:

	income	happiness
0	3.862647	2.314489
1	4.979381	3.433490
2	4.923957	4.599373
3	3.214372	2.791114
4	7.196409	5.596398

In [5]: `y = df['happiness']`
`x = df['income']`

In [6]: **class SimpleLinearRegression:**

```

def __init__(self):
    """
    Initializes the SimpleLinearRegression model.

    Attributes:
    -----
    B_0 : float
        The intercept of the regression line.
    B_1 : float
        The slope of the regression line.
    MSE : float
        Mean squared error.
    r_squared : float
        Coefficient of determination.
    """
    self.B_0 = None
    self.B_1 = None
    self.MSE = None
    self.r_squared = None

def fit(self, X, y):
    """
    Fits the simple linear regression model to the training data. Calculates and sets the intercept (B_0) and slope (B_1) of the regression line and then calculates the evaluation metrics for the model which are MSE and r_squared.

    Parameters:
    -----
    X : array-like
        The input feature array.
    y : array-like
        The target variable array.
    """
    self.n = len(X)
    # Ensure X and y are numpy arrays
    self.X = np.array(X)
    self.y = np.array(y)

    # Calculate the mean of X and y
    self.x_bar = X.mean()
    y_bar = y.mean()

    # Calculate SXX and SXY
    self.SXX = sum(X**2) - self.n * self.x_bar**2
    SXY = sum(X * y) - self.n * self.x_bar * y_bar

    # Calculate coefficients
    self.B_1 = SXY / self.SXX
    self.B_0 = y_bar - self.B_1 * self.x_bar

    # Calculate estimated error
    y_hat = self.predict(X)
    e = y - y_hat
    self.SSE = sum(e**2)
    self.MSE = self.SSE / (self.n - 2)

```

```

# Calculate coefficient of determination
self.SST = sum((y - y_bar)**2)
self.SSR = sum((y_hat - y_bar)**2)
self.r_squared = self.SSR / self.SST

def predict(self, X):
    """
    Predicts the target variable for new input values using the fitted
    model.

    Parameters:
    -----
    X : array-like
        The input feature array for which to predict the target variable.

    Returns:
    -----
    array-like
        Predicted values of the target variable.

    Raises:
    -----
    ValueError
        If the model has not been fitted yet.
    """
    if self.B_0 is None or self.B_1 is None:
        raise ValueError("The model has not been fitted yet.")

    X = np.array(X)
    return self.B_0 + self.B_1 * X

def plot(self, X, y):
    """
    Plots the scatter plot of the data points and the regression line.

    Parameters:
    -----
    X : array-like
        The input feature array.
    y : array-like
        The target variable array.
    """
    if self.B_0 is None or self.B_1 is None:
        raise ValueError("The model has not been fitted yet.")

    # Ensure X and y are numpy arrays
    X = np.array(X)
    y = np.array(y)

    # Predicted values
    y_pred = self.predict(X)

    # Plot
    plt.scatter(X, y, color='blue', label='Data points')
    plt.plot(X, y_pred, color='red', label='Regression line')
    plt.xlabel('X')
    plt.ylabel('y')
    plt.title(f'Simple Linear Regression with line of best fit: y = {round(self.B_0,2)} + {round(self.B_1,2)} x')
    plt.legend()

```

```

plt.show()

def anova_table(self):
    """
    Constructs and returns the ANOVA table.

    Returns:
    -----
    pd.DataFrame
        The ANOVA table showing the Sum of Squares, Degrees of Freedom,
Mean Squares, F-Statistic,
        and p-value for the model.
    """
    if self.B_0 is None or self.B_1 is None:
        raise ValueError("The model has not been fitted yet.")

    # Degrees of freedom
    self.df_regression = 1
    self.df_error = self.n - 2
    df_total = self.n - 1

    # Mean squares
    MSR = self.SSR / self.df_regression
    MSE = self.SSE / self.df_error

    # F-statistic
    self.F_stat = MSR / MSE

    # Assemble the ANOVA table
    anova_data = {
        'Source': ['Regression', 'Error', 'Total'],
        'Sum of Squares': [self.SSR, self.SSE, self.SST],
        'Degrees of Freedom': [self.df_regression, self.df_error, df_to
tal],
        'Mean Square': [MSR, MSE, ""],
        'F-Statistic': [self.F_stat, "", ""]
    }

    anova_table = pd.DataFrame(anova_data)
    return anova_table

def hypothesis_test(self, alpha=0.05):
    """
    Perform a hypothesis test to determine the relationship between the
independent variable (x)
    and the dependent variable (y) based on the F-statistic.

    Parameters:
    -----
    alpha : float
        The significance level for the test. Default value = 0.05

    Returns:
    -----
    None
        The function prints the hypothesis test results, including the
F-statistic,
        critical value, and the conclusion.

    Notes:

```

```

-----
    The null hypothesis ( $H_0$ ) assumes that there is no relationship between  $x$  and  $y$ ,
    while the alternative hypothesis ( $H_a$ ) suggests that there is a relationship.
    """
    # Calculate the critical value for the given significance level
    F_c = scipy.stats.f.ppf(1 - alpha, self.df_regression, self.df_error)

    # Determine the conclusion based on the F-statistic
    if self.F_stat > F_c:
        conclusion = ("Since  $F_0 > F_c$ , we reject the null hypothesis.
        Therefore, there's a relationship between  $x$  and  $y$ ."
        )
    else:
        conclusion = ("Since  $F_c > F_0$ , we don't reject the null hypothesis.
        Therefore, there's no relationship between  $x$  and  $y$ ."
        )

    # Print the results in a structured and visually appealing format
    print("="*45)
    print("          Hypothesis Testing Results          ")
    print("="*45)
    print(f"{'Null Hypothesis ( $H_0$ ):':<25}  $B_1 = 0$ ")
    print(f"{'Alternative Hypothesis ( $H_a$ ):':<25}  $B_1 \neq 0$ ")
    print("-"*45)
    print(f"{'F-statistic ( $F_0$ ):':<25} {self.F_stat:.4f}")
    print(f"{'Critical value ( $F_c$ ):':<25} {F_c:.4f}")
    print("-"*45)
    print(f"{conclusion}")
    print("="*45)

    def interval_estimation(self, alpha=0.05, sigma=None):
        if sigma is None:
            t = scipy.stats.t.ppf(1-(alpha/2), self.df_error)
            B_0_lower = self.B_0 - t * np.sqrt(self.MSE*((1/self.n) + (self.x_bar/self.SXX)))
            B_0_upper = self.B_0 + t * np.sqrt(self.MSE*((1/self.n) + (self.x_bar/self.SXX)))
            B_1_lower = self.B_1 - t * np.sqrt(self.MSE/self.SXX)
            B_1_upper = self.B_1 + t * np.sqrt(self.MSE/self.SXX)
        else:
            z = scipy.stats.norm.ppf(1-(alpha/2))
            B_0_lower = self.B_0 - z * sigma * np.sqrt((1/self.n) + (self.x_bar/self.SXX))
            B_0_upper = self.B_0 + z * sigma * np.sqrt((1/self.n) + (self.x_bar/self.SXX))
            B_1_lower = self.B_1 - z * sigma * np.sqrt(1/self.SXX)
            B_1_upper = self.B_1 + z * sigma * np.sqrt(1/self.SXX)

        return ([B_0_lower, B_0_upper],[B_1_lower, B_1_upper])

```

```

In [7]: model = SimpleLinearRegression()
        model.fit(x, y)

```

```
In [8]: print("Evaluation of the model")
print("-----")
print(f'Line of best fit is: y = {round(model.B_0,2)} + {round(model.B_1,2)} x')
print(f'Mean squared error is: {round(model.MSE,3)}')
print(f'Coefficient of determination is: {round(model.r_squared,3)}')
```

Evaluation of the model

Line of best fit is: $y = 0.2 + 0.71 x$

Mean squared error is: 0.516

Coefficient of determination is: 0.749

```
In [9]: model.anova_table()
```

Out[9]:

	Source	Sum of Squares	Degrees of Freedom	Mean Square	F-Statistic
0	Regression	764.546359	1	764.546359	1482.632007
1	Error	255.771488	496	0.515668	
2	Total	1020.317847	497		

```
In [10]: model.hypothesis_test()
```

```
=====
Hypothesis Testing Results
=====
Null Hypothesis (H_0): B_1 = 0
Alternative Hypothesis (H_a): B_1 ≠ 0
-----
F-statistic (F_0): 1482.6320
Critical value (F_c): 3.8603
-----
Since F_0 > F_c, we reject the null hypothesis.
Therefore, there's a relationship between x and y.
=====
```

```
In [11]: B_0_est, B_1_est = model.interval_estimation()
print(f"95% CI for B_0: {B_0_est}")
print(f"95% CI for B_1: {B_1_est}")
```

95% CI for B_0: [0.1046540279959846, 0.3038867644124284]

95% CI for B_1: [0.6774017746996904, 0.7502492498607113]

```
In [12]: B_0_est, B_1_est = model.interval_estimation(sigma=x.std())
print(f"95% CI for B_0: {B_0_est}")
print(f"95% CI for B_1: {B_1_est}")
```

95% CI for B_0: [-0.03617479319616132, 0.44471558560457436]

95% CI for B_1: [0.6259091122071826, 0.8017419123532191]

```
In [ ]:
```